

Christopher Saw, Jose Goodenough-Shaw

ECS 188

June 5, 2025

*Automating Inequality: How High-Tech Tools Profile, Police, and Punish the Poor*

At the core of a lot of the wave of automation is an appeal to impartiality of machines. For how can code be wrong? Code is just an unflinchingly unfailing collection of rules. Rules that a machine follows to a fault. But code is not some altruistic thing that appears out of nowhere. It is shaped by human hands and human hands are ultimately flawed.

When looking at automated systems as those examined by Eubanks there is often an appeal to brutal honesty. That while it is a shame that these negative outcomes are an unfortunate result, nevertheless this is the best way. The best utilization of resources. Eubanks explores this logic through three case studies: Indiana's automation of welfare benefits, Los Angeles' Coordinated Entry System for housing the unhoused, and Allegheny County's predictive child welfare risk model. Each system, while promising efficiency, ultimately reinforces structural inequality.

A common argument from proponents of the coming age of automation is one of utility. These advancements are framed from the lens of efficiency and production. While some rhetoric may involve how humans benefit, overwhelmingly the focus is on the beauty of efficiency. Yet when looking at the actual original edicts of utilitarianism itself, one cannot fully divorce the human element these industrialists are so eager to do. While there is much debate on how to define happiness and utility, Utilitarianism strongly declares that what benefits humanity as a whole to be of the utmost importance. Looking at the automation of welfare eligibility in Indiana we can see while there is an argument towards utility, it is not one based in benefiting the

majority of people. Quoting the then governor of Indiana at the time Mitch Daniels one can see a perfect exemplar of the kind of rhetoric being referenced ,

" 'irretrievably broken,' wasteful, fraudulent, and "America's worst welfare system."... In early 2006, the Daniels administration released a request for proposal (RFP) to outsource and automate eligibility processes for TANF, food stamps, and Medicaid. In the request, the state set very clear goals: reduce fraud, curtail spending, and move clients off the welfare rolls." (Eubanks 44).

Daniels and others like him are right in some ways. Human systems can be just as flawed if not more so than algorithmic ones. Eubanks herself cites several caseworkers who attest to the inefficiencies inherent to the welfare system as it was. Managing the needs of millions is an incredibly difficult challenge and automated systems could be a boon in aiding and streamlining the process. Where they fail is in their definition of utility. State actors which make sweeping decisions like this do not attempt to quantify happiness or do hedonic calculus. While public opinion may be polled it rarely is the end all be all of any action such an actor takes. No, what utility means now more than ever in the modern age is capital. The rhetoric of inefficiency is the inefficiency of tax dollars, the monetary waste of supporting the poor. In the wake of automation more than a million applications for welfare were rejected in Indiana, a figure Eubanks notes as a 54% increase from before automation. But even beyond those figures Eubanks paints a picture of widespread suffering from those who did manage to get their applications accepted. Endless forms and phone calls and bureaucratic busywork colors the experience of those trying to apply for welfare (Eubanks 50). While these automations are framed as all powerful machines that vastly improve productivity, perhaps the Mechanical Turk is a better analogy. The systems aren't really more efficient with automation. They have simply offloaded the work onto the poor and

less fortunate. If we then approach from the Greatest Happiness Principle, one could argue the net change in happiness is negligible or even negative. After all the same amount of work is being done, just by different people. Not to mention the vast increase in negative outcomes. At best one could argue that perhaps the money saved could then be spent on other programs that could benefit more people, but going back to that original rhetoric there is no suggestion of such programs. Rather, it's all focussed on cutting chaff. The goals were not to improve the lives of people, it was to "reduce fraud, curtail spending, and move clients off the welfare rolls." (Eubanks 44).

When harm becomes routine and predictable, the question is no longer whether a system works—but whether it deserves to exist at all. One of the most fundamental ethical problems with automated public service systems is their lack of legitimacy and meaningful consideration for the people they affect. At first glance, one might argue that these systems possess a degree of democratic legitimacy since they are often implemented through legislation, approved by elected officials, and publicly framed as commonsense solutions to inefficiency and waste. Their continued existence, some would claim, is evidence that they have passed through the appropriate channels of consent. But this view assumes that consent in a bureaucratic democracy is meaningful, when in fact it can be deeply limited—especially for those whose survival depends on systems they cannot afford to question. To truly evaluate the legitimacy of automated welfare systems, we must apply a more rigorous ethical standard. John Rawls' "veil of ignorance" thought experiment asks us to design social institutions as if we had no knowledge of our class, race, or position in society (Graham 2010, p. 18). From this perspective, no rational person would agree to live in a society where access to life-sustaining resources—such as housing, food or the privilege to care for your own child—is determined by blurred algorithms that offer no

transparency, accountability or clear process for appeal. A system that might seem acceptable when one imagines themselves as a decision-maker or taxpayer begins to look far less just when one imagines being on the receiving end of its failure. Eubanks' case study in Allegheny County further highlights how deeply this problem runs. There, families are scored by a predictive risk model designed to forecast child abuse or neglect. Many of these families never even know they've been "assessed", let alone receive a clear explanation of the scoring process or its consequences (Eubanks, 2018, p. 103). In essence, their fates are sealed in a data warehouse. This is not democratic governance-it's surveillance disguised as support. According to Rawls, inequalities are only ethically justifiable if they improve the condition of the least advantaged members of society (Graham, 2010, p. 18).. But what Eubanks shows, again and again, is that automated welfare systems are built not to elevate the vulnerable but to regulate, punish and push them further to the margins. This contradiction becomes especially obvious in Eubank's analysis of Indiana, where automated eligibility systems resulted in over a million welfare denials-illustrating how these tools often generate exclusion rather than support. The official justification was cost saving and fraud reduction-but the real outcome was mass exclusion (Eubanks, 2018, pp. 143). Elderly, disabled and impoverished residents were denied basic aid, often for minor bureaucratic technicalities interpreted by machines as "failure to cooperate" (Eubanks, 2018, p. 17). They had no meaningful way to challenge or even to begin to understand why they were rejected aid. Proponents of these systems often invoke the idea of tacit consent: that participation in a system implies agreement with its structure. But as Graham argues, "Tacit consent is valid only if dissent is genuinely possible" (Graham, 2010, p. 49). When survival depends on cooperation with an unaccountable system, compliance isn't consent-it's coercion. These systems do not extend the social contract; they break it. What emerges is not a new form

of ethical governance but a technological sleight of hand that hides oppression behind efficiency. Far from embodying Rawlsian justice, these systems represent a betrayal of it-replacing public reason with code, and fairness with forced silence.

So if these systems are both negative and illegitimate, who is responsible for their actions? When speaking of philosophy a human actor is presupposed. As Korsgaard puts it, “[human beings] are condemned to choice and action” (Korsgaard 1). And in most of the philosophies covered by Graham only humans are considered valid actors in their system. Yet in this age of technology machines are often either making the final decision or greatly influencing them. Proponents of these technologies see this as a positive, appealing to high minded ideas about innovation and a new digital age. Eubanks notes this phenomenon herself stating,

“When we talk about the technologies that mediate our interactions with public agencies today, we tend to focus on their innovative qualities, the ways they break with convention. Their biggest fans call them “disruptors,” arguing that they shake up old relations of power, producing government that is more transparent, responsive, efficient, even inherently more democratic.” (Eubanks 176)

Yet again there is that divorce from the machine and the human. These systems are spoken separately from their creators, as if they are a wholly independent object. But as Eubanks argues, it misses how these systems are inevitably intertwined with their creators and their own historical baggage. Throughout the book in every one of these case studies Eubanks connects these systems with those in the past. The overarching historical system she constantly refers back to are the poorhouses of the 19th century, drawing parallels between them in how they demean and hurt the poor. And when these developers, these State actors simply implement these systems and

try to wash their hands of it, try to claim that any negative decisions are just the brutal honesty of a machine, well it reeks of bad faith. A way to dodge and abdicate responsibility.

Even beyond the question of should is one of capability While it may be easy to say a machine is perfectly capable with making decisions for minor things like checking the fulfillment of boxes on a form, the question of whether a machine can decide become vastly more complex when looking at the cases that Eubanks examines as in with the Allegheny Family Screening Tool (or AFST). Even prior to AFST being implemented Eubanks notes how much grey area there is in what the government considers abuse and neglect, with the final decision being left up to the discretion of case workers. A majority of cases are focussed on neglect and Eubanks notes how the conditions of neglect and poverty are incredibly blurred (Eubanks 129). When testing the algorithm herself against a caseworker's assessment of a case an immediate disconnect arises. Where the caseworker rated the risk of child harm for the test case to be a 6 out of 20, the algorithm decided it was a 14 (Eubanks 129). There is an opaqueness to the algorithm's work as compared to the caseworker's. Where the caseworker makes an argument, cites specific rules and precedent, the algorithm just spits out a number. A possible huge life changing decision to separate a child from their family is reduced to a simple number. In its current form the algorithm is intended to merely be a tool for caseworkers. Another system through which they can make their final judgement. However Eubanks notes that while that may be the intent, that is not how it works in practice,

“We all tend to defer to machines, which can seem more neutral, more objective. But it is troubling that managers believe that if the intake screener and the computer's assessments conflict, the human should learn from the model. The AFST, like all risk models, offers only probabilities, not perfect prediction” (Eubanks 141).

These machines are not perfect. Eubanks goes into further exhaustive detail on how algorithms like the AFST have many failures and are really only a predictor rather than a determiner. Yet again and again they are relied on and used to influence incredibly important decisions with no form of accountability.

Some may argue still that while there are negative outcomes and failures in these systems now, it is simply an issue of technology. That in time these issues will be smoothed over by the ingenuity of software engineers. But that fails to see the underlying issue at hand. If left unchecked, the logic of algorithmic governance can spiral into a quiet form of eugenics—a possibility Virginia Eubanks does not shy away from in *Automating Inequality*. She draws parallels between modern welfare technologies and the historical project of eugenics, where data and scientific authority were weaponized to categorize lives as “fit” or “unfit” (Eubanks, 2018, p. 24). Today’s predictive risk models—like those used in Los Angeles’ Coordinated Entry System or Allegheny County’s child welfare screening—don’t assess individuals based on present behavior. Instead, they sort people by statistical proxies, often derived from historically biased data. These systems don’t just anticipate harm; they effectively assign blame, punishing people for poverty, race, zip code, or other inherited disadvantages. This shift reflects more than flawed data—it reveals a broader ideological failure. The just-world fallacy—the belief that people get what they deserve—offers moral cover for such systems. It allows those in power to view their own privilege as earned and to regard the struggles of the poor as self-inflicted. When these beliefs are encoded into algorithms, they become operationalized prejudice. The result is a system that quietly disciplines the poor, discouraging reproduction, monitoring for deviance, and withdrawing aid in the name of efficiency. As Gordon Graham puts it, “Utilitarianism... has always run the risk of sacrificing the interests of individuals in the name of maximizing the

happiness of the majority.” (Graham, 2010, p. 102). That tradeoff, once theoretical, now unfolds in code—justified by spreadsheets, cloaked in objectivity, and executed at scale. Utility becomes the rationale for exclusion, and data becomes the mask that hides it. The focus shifts from improving lives to managing scarcity—efficiently, impersonally, and without moral accountability. If justice is to remain a goal of public systems, we must resist the belief that ethical governance can be automated. These systems do not merely reflect inequality—they institutionalize it. A just society cannot measure human worth in probabilities of risk, nor reduce the lives of the vulnerable to variables in a cost-benefit analysis.

In the end what Eubanks sees as the issues with these systems is not the presence of technology itself but of people. These systems cannot be divorced from the culture that birthed them, and unfortunately that culture is one that vilifies poverty. Effective change will not come solely from technological advancements, but rather widespread social change. Eubanks describes the issue of poverty as,

“... a borderland. There’s quite a lot of movement in the economic fringes, especially across the fuzzy boundary between the poor and the working class. Those who live in the economic borderlands are pitted against one another by policies that squeeze every possible dime from the wallets of the working class at the same time that they cut social programs for the poor and absolve the professional middle class and wealthy of their social obligations. There is a lot of self-blame and horizontal violence in the borderlands, but there is also a lot of shared experience. The first challenge we face in dismantling the digital poorhouse is building empathy and understanding among poor and working-class people in order to forge winning political coalitions.” (Eubanks 205).



These new algorithms are not the cause, they are simply the newest tool of enforcement. When discussing the role of technology in this dismantling, Eubanks pointedly refers to two important questions, “Does the tool increase the self-determination and agency of the poor? Would the tool be tolerated if it was targeted at non-poor people?” (Eubanks 211). It brings to mind again that idea of the veil of ignorance. Perhaps it is the most fitting framework to end off with. The demonization of the poor largely comes from a perspective of moral superiority, while the framework of the veil acknowledges how being poor is a very real possibility for the majority of people.

## Works Referenced

Eubanks, Virginia. *Automating Inequality: How High-Tech Tools Profile, Police, and Punish the Poor*. Picador, St. Martin's Press, 2019.

Graham, Gordon. *Theories of Ethics: An Introduction to Moral Philosophy with a Selection of Classic Readings*. Routledge : Imprint of the Taylor & Francis Group, 2011.

Korsgaard. *Self-Constitution: Agency, Identity, and Integrity*. Oxford University Press, 2009.